

REYNALD ACE PILPIL

0968 619 1930 | azepilpil139@gmail.com | reynaldace-portfolio.online | github.com/aces-14 | huggingface.co/acepilpil

EDUCATION

Technological University of the Philippines – Cavite

2021 – 2026

Bachelor of Engineering Technology — Major in Computer Engineering

CERTIFICATIONS

Google Data Analytics Professional Certificate • Coursera (Google)

July 2024

Introduction to LangGraph (Python) • LangChain Academy

May 2026

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript

Machine Learning & AI: Machine Learning, Deep Learning, NLP, LLMs, Computer Vision, Predictive Modeling, Feature Engineering, Model Evaluation, Explainable AI, PyTorch, scikit-learn, XGBoost, Hugging Face Transformers, pandas, NumPy

LLM & Agentic Systems: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Multi-Agent Systems, Vector Databases, Prompt Engineering, LangChain, LangGraph, LangSmith

APIs & Backend: REST APIs, API Design, FastAPI, Django, Flask, Docker, MySQL, SQL Optimization

MLOps & Deployment: Cloud Deployment, Containerization, CI/CD, Streamlit, Gradio, Hugging Face Spaces, Railway

Tools & Practices: Git, GitHub, Version Control, Agile, Code Review, Testing, Debugging, Power BI

EXPERIENCE

Technical Sales Intern • 3E HITECH SOLUTIONS, INC.

January 2026 – May 2026

- Supported pre-sales activities focused on projects with the National Grid Corporation of the Philippines (NGCP) — contributing to proposal preparation, scoping, and technical clarifications with stakeholders.
- Assisted project management across NGCP deployments — coordinating timelines, tracking deliverables, and supporting communication between technical teams and clients.

KEY AI / MACHINE LEARNING PROJECTS

SentinelPH • Dengue Risk Intelligence Platform

Tech: Streamlit, XGBoost, LangGraph, LangChain, ChromaDB, Groq (Llama 3.3 70B), HuggingFace · **Live:** [Portfolio](#) · [GitHub](#)

- Architected an end-to-end dengue risk dashboard for the Philippines combining four data sources (OpenDengue 2012–2023, Open-Meteo weather, Google Trends, GDELT news) into a cleaned, merged Parquet pipeline feeding the model and chatbot.
- Trained an XGBoost forecast model on lag features, rolling means, weather lags, and seasonal flags to predict next-week case counts and risk levels (LOW/MEDIUM/HIGH) — 40% more accurate than a naive seasonal baseline.
- Implemented a 4-agent LangGraph briefing system with an evaluator agent that independently audits each report against data grounding, uncertainty language, and medical-claim avoidance — rewriting on failure.

Market Intelligence Agent • Stateful Multi-Agent System with Self-Revision

Tech: Python, LangGraph, LangChain, Groq (Llama 3.3 70B), Tavily Search API, Pydantic v2, Gradio, LangSmith · **Live:** [HF Space](#) · [GitHub](#)

- Built an autonomous research agent with a self-evaluating revision loop using LangGraph — agent critiques its own draft, scores it, and re-searches targeted gaps until a quality threshold is met.
- Researches a company across four parallel angles (overview, news, competitors, market trends) via live web search and synthesizes results into a structured Word document.

Job Application Assistant Agent • Autonomous Multi-Step LLM Agent

Tech: Python, Groq (Llama 3.3 70B), Tavily Search API, Pydantic v2, Gradio, python-docx · **Live:** [HF Space](#) · [GitHub](#)

- Built a 9-step agent with integrity enforcement — architecturally prohibited from fabricating resume content; post-tailoring diff check flags violations with an auditable changelog (reason + confidence score).
- Pipeline: resume parsing → JD scoring → company research via live web search → ATS-tailored resume, cover letter, and LinkedIn outreach delivered as Word documents in under 30 seconds.

Philippine Labor Law RAG Assistant • End-to-End Retrieval-Augmented Generation

Tech: Python, LangChain, ChromaDB, Hugging Face Transformers, Flan-T5, sentence-transformers, Gradio · **Live:** [HF Space](#)

- Designed and deployed an end-to-end RAG pipeline over the Philippine Labor Code — document ingestion, semantic chunking, vector embeddings, ChromaDB retrieval, Flan-T5 generation.
- Returns natural-language answers (leave entitlements, overtime, termination procedures) with source citations grounded in the actual law, reducing hallucination through retrieval-first design.