

REYNALD ACE PILPIL

Cavite, Philippines | 0968 619 1930 | azepilpil139@gmail.com
reynaldace-portfolio.online | github.com/aces-14 | huggingface.co/acepilpil

EDUCATION

Technological University of the Philippines – Cavite

Bachelor of Engineering Technology, Major in Computer Engineering

CERTIFICATIONS

Google Data Analytics Professional Certificate — Coursera (Google)

July 2024

Introduction to LangGraph (Python) — LangChain Academy

May 2026

TECHNICAL SKILLS

Languages & Tools: Python, SQL, Git/GitHub

Machine Learning: Supervised & unsupervised learning, feature engineering, model training & evaluation, explainable AI, scikit-learn, gradient boosting (XGBoost), classification & regression models

Deep Learning & NLP: Neural networks, Transformers, Natural Language Processing (NLP), fine-tuning, embeddings, document AI, PyTorch, Hugging Face

LLMs & Agents: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), AI agents & multi-agent systems, prompt engineering, vector databases, LangChain, LangGraph

MLOps & Deployment: FastAPI, REST APIs, Docker, model deployment, cloud deployment

EXPERIENCE

Technical Sales Intern — 3E HITECH Solutions, Inc.

January 2026 – May 2026

- Supported pre-sales for National Grid Corporation of the Philippines (NGCP) projects — proposal preparation, scoping, and technical clarifications with stakeholders.
- Assisted project management across NGCP deployments — coordinating timelines, tracking deliverables, and supporting team-to-client communication.

KEY AI / MACHINE LEARNING PROJECTS

SentinelPH — Dengue Risk Intelligence Platform

Tech: Streamlit, XGBoost, LangGraph, LangChain, ChromaDB, Groq (Llama 3.3 70B), Hugging Face · *Live:* [Portfolio](#) · [GitHub](#)

- Architected a dengue risk dashboard merging four data sources (OpenDengue, Open-Meteo, Google Trends, GDELT) into a cleaned Parquet pipeline feeding the model and chatbot.
- Trained an XGBoost model on lag features, rolling means, and seasonal flags to forecast next-week cases and risk levels (LOW/MEDIUM/HIGH) — 40% more accurate than a naive seasonal baseline.
- Built a four-agent LangGraph briefing system with an evaluator agent that audits each report for grounding, uncertainty, and medical-claim avoidance, rewriting on failure.

Market Intelligence Agent — Stateful Multi-Agent System with Self-Revision

Tech: Python, LangGraph, LangChain, Groq (Llama 3.3 70B), Tavily Search API, Pydantic v2, Gradio, LangSmith · *Live:* [HF Space](#) · [GitHub](#)

- Built an autonomous research agent with a self-evaluating revision loop in LangGraph; it critiques its own draft, scores it, and re-searches gaps until a quality threshold is met.
- Researches a company across four angles (overview, news, competitors, market trends) via live web search, synthesizing results into a structured Word document.

Job Application Assistant Agent — Autonomous Multi-Step LLM Agent

Tech: Python, Groq (Llama 3.3 70B), Tavily Search API, Pydantic v2, Gradio, python-docx · *Live:* [HF Space](#) · [GitHub](#)

- Built a nine-step agent with integrity enforcement — prohibited from fabricating resume content; a post-tailoring diff check flags violations with an auditable changelog (reason + confidence).
- Designed a parse-to-delivery pipeline that ingests a resume, scores it against the job description, researches the company, and outputs an ATS-tailored resume, cover letter, and LinkedIn message in under 30 seconds.

Philippine Labor Law RAG Assistant — End-to-End Retrieval-Augmented Generation

Tech: Python, LangChain, ChromaDB, Hugging Face Transformers, Flan-T5, sentence-transformers, Gradio · *Live:* [HF Space](#)

- Designed and deployed a RAG pipeline over the Philippine Labor Code — document ingestion, semantic chunking, vector embeddings, ChromaDB retrieval, and Flan-T5 generation.
- Returns cited natural-language answers (leave, overtime, termination) grounded in the actual law, reducing hallucination through retrieval-first design.